

## Programming Assignment 8

### Due Sunday, April 5

#### Objectives

Set up and test a GPU environment. Write and test a simple GPU program.

#### Prerequisites (covered in class or through examples in on-line documentation)

**HW** – GPU basics: accelerators, accelerator memory, SIMT support

**SW** – SIMT model, basic CUDA model

**Programming** – CUDA environment, CUDA/GPU set up and data transfer, simple GPU kernels

---

#### Assignment

##### Part 1: Tutorial

Follow the directions in `cuda_example-2.doc` to set up your environment. Read `cuta_test.cu` and make sure you understand it. Compile and run the program.

##### Part 2: A simplified SOR

1. Create a simplified SOR based on the sample code given in class (for the 1D averaging filter). It should do the following:

- On the host: Create and initialize a 2Kx2K array of single precision floats
  - Transfer the array to the GPU
  - Create a single block of 16x16 threads (extend the 1D example from the slides to 2D)
  - Let each thread operate on a square patch of the input array (rather than a single output element)
  - Run 2000 iterations of SOR – the boundary should not change. You do not need to test for convergence.
  - Transfer the output back to the CPU
2. Create a reference code that performs the same function on the host. Test your answer against the same procedure on the host. Are they different? How?
3. Time both codes (GPU and CPU). How long do they take?
4. (Optional) Threads are responsible for multiple elements. Try different assignments. For example: 1D strip per thread, 2D tile per thread, interleaved (every 256<sup>th</sup> element).

##### Part 3: Multiple blocks

Make the following changes to your GPU code. Check your answer and your timing.

1. Let each thread operate on a single output array element (like the code in class)
2. Let the block size be 16x16. Choose the grid dimensions accordingly.
3. Modify the host code (kernel calls) so that the program executes correctly. That is, the iterations should now be controlled by the host not the GPU (why?).
4. Check correctness and timing.
5. (Optional) Try using thread control as in Part 2. For this to be interesting should also try different mappings of array elements to threads.

NOTE: Because of the difference in precision and various floating point technical issues, you will probably see some drift between the GPU version and the reference version even with correct code. You should be able to handle this by increasing the epsilon to, say, 5%.

NOTE: Last year many students found `cuPrintf` useful; the relevant files are attached as part of this assignment.