

Review: Concurrent Processes

Outline

1. Review basic issues in concurrency
2. The critical section problem
3. Hardware support for critical sections

Concurrent Processes

Consider the following program:

S1: **a** := **x** + **y**;

S2: **b** := **z** + **1**;

S3: **c** := **a** - **b**;

S4: **w** := **c** + **1**;

Which can be executed concurrently?

Which must precede others?

Let each node be a process


Then each FORK initiates a new process

Each JOIN kills all but one process

S1 must precede S3

S2 must precede S3

S3 must precede S4



How to Determine Precedence

Fundamental Property: If statements are not explicitly constrained, then any order is possible. >> *would reordering change the outcome?*

→ True for serial as well as parallel systems!

Q: If an instruction precedes another, when can this precedence be removed without changing the intended result of the program execution?

A: When three conditions hold:

a) No Read after Write (data/true dependence)

```
a := b + c;
```

```
d := a + e;
```

b) No Write after Read (anti-dependence)

```
d := a + e;
```


```
a := b + c;
```

c) No Write after Write (output dependence)

```
a := b + c;
```

```
a := d + e;
```

Precedence Graphs




Arrows represent orderings
→ arrows are transitive
→ graphs must be acyclic

Handmade precedence graphs

forkL \equiv produces two concurrent executions
1) continuation after fork
2) execution beginning at label L

Example:


S1 ;
fork L ;
S2 ;
.
.
L: S3 ;



join count \equiv recombine multiple concurrent threads into one


Example 2:

```
count := 2;  
fork L1;  
a := x + y;  
goto L2;  
L1: b := z + 1;  
L2: join count  
c := a - b;  
w := c + 1;
```



Example 3:

```
S1;  
count := 3;  
fork L1;  
S2;  
S4;  
fork L2;  
S5;  
goto L3;  
L2: S6;  
goto L3;  
L1: S3;  
L3: join count;  
S7;
```




Problem: Too much flexibility -- makes efficient compilation and code maintenance very difficult

Specification Alternative: Parbegin/Parend

S0;
Parbegin S1; S2; ... Sn; Parend
Sn+1;

≡ All statements S_i between
Parbegin and Parend can be
executed simultaneously



Original example

Parbegin


a := x + y; # S1


b := z + 1; # S2

Parend

c := a - b; # S3

w := c + 1; # S4






```

# 2nd Example
# Note block
# structured style
S1;
Parbegin
  S3;
  Begin
    S2;
    S4;
    Parbegin
      S5;
      S6;
    Parend
  End
Parend
S7;
  
```

Another example:
easy with fork/join but impossible to
implement with parbegin/parend



The Critical Section Problem

Let P1 and P2 be two processes ...

Let Count = 5

```
Begin P1
```

```
·
```

```
·
```

```
Count := Count + 1;
```

```
·
```

```
End P1
```

```
Begin P2
```

```
·
```

```
Count := Count - 1;
```

```
·
```

```
·
```

```
End P2
```

What is the value of Count when P1 and P2 finish?

Begin P1

```
.  
.   
  Count := Count + 1;  
.
```

End P1

in assembly language

```
LD  R1,Count  
ADD R1,R1,1  
ST  R1,Count
```

Begin P2

```
.   
  Count := Count - 1;  
.
```

End P2

in assembly language

```
LD  R2,Count  
SUB R2,R2,1  
ST  R2,Count
```

```
# One possible interleaving:  
# (Why are multiple different interleavings  
#   possible even on a single core?)
```

```
LD  R1,Count      # P1  
ADD R1,R1,1       # P1  
LD  R2,Count      # P2  
SUB R2,R1,1       # P2  
ST  R2,Count      # P2  
ST  R1,Count      # P1
```

```
# PROBLEM:  Count must be updated ATOMICALLY!
```

Solution Requirements

3 Issues

1. Mutual Exclusion: If P_i is in the critical section, then no other P_j is allowed in.
2. Progress: Only processes waiting to enter the critical section can participate in the decision. Decision must be made in finite time.
3. Bounded Waiting: All processes are allowed in the critical section only a finite # of times before every process has access.

Critical Section Problem: SW Solution

```
while ( ) do
  begin
    parbegin
      P0;
      P1;
    parend
  end
end
```

Process P_i

repeat

entry section

critical section

exit section

remainder section

until false

Algorithm 1

```
BOOL TURN = RANDOM();  
// Two processes: P0, P1  
FOR Pi Initialize TURN to 0 or 1  
REPEAT  
    entry section  
    WHILE TURN  $\neq$  i DO SKIP;  
  
    critical section  
  
    TURN := j;  
    remainder section  
UNTIL FALSE;
```

Allows alternating execution
What's the problem?

Algorithm 2

```
// Two processes: P0, P1
flag[i] := FALSE; flag[j] := FALSE;
REPEAT
    entry section
    WHILE flag[j] DO SKIP;
    flag[i] := TRUE;

    critical section

    flag[i] := FALSE;
    remainder section
UNTIL FALSE;
```

What's the problem?

T0: P0 executes `while` and finds `flag[1] == F`

T1: P1 executes `while` and finds `flag[0] == F`

T2: P1 sets `flag[1] ← T`

T3: P0 sets `flag[0] ← T`

Now BOTH P0, P1 are in the critical section!

Q: Can we make

*while ...
flag ...*

another critical section?

Algorithm 3

```
// Two processes: P0, P1
REPEAT
    entry section
    flag[i] := TRUE;
    WHILE flag[j] DO SKIP;

    critical section

    flag[i] := FALSE;
    remainder section
UNTIL FALSE;
```

What's the problem?

Solution

```
// Two processes: P0, P1
flag[0], flag[1]; // initialize to FALSE
turn; // initialize to 0 or 1
REPEAT
    entry section
    flag[i] := TRUE; // I'd like to enter
    turn := j; // but it's yours if you want it
    // I'll wait while you want it or it's your turn
    WHILE (flag[j] AND turn = j) DO SKIP;

    critical section

    flag[i] := FALSE; // I don't want it anymore
    remainder section
UNTIL FALSE;
```


Why does it work?

→ setting `flag[i] = FALSE` when leaving the critical section ensures that this process won't prevent entry

→ `turn` ensures that only one process will succeed

a) either it is P_i 's turn or P_j does not care

b) c) only loop if `flag[j] = TRUE AND turn = j`

This only happens if P_j is in the critical section

What about crash in the critical section?
... mendacity??

Hardware Solutions

Basic Technique: Multiple operations are performed on one memory location ***atomically***.

Two Methods:

```
int TEST-AND-SET(target)
    TEST-AND-SET := target;
    target := TRUE
```

```
VOID SWAP(a,b)
    temp := a;
    a := b;
    b := temp;
```

Critical Section w/ atomic primitives

```
// Use TEST-AND-SET
// Let "target" be lock
// Let lock be FALSE
REPEAT
    entry section
    WHILE TEST-AND-SET(lock)
        DO SKIP;

    critical section

    lock := FALSE;
    remainder section
UNTIL FALSE;
```

```
// Use SWAP
// Let "lock" be a global Boolean
// Let "key" be a local Boolean
// Initialize lock to FALSE
REPEAT
    entry section
    key := TRUE;
    REPEAT
        SWAP(lock, key);
    UNTIL key = FALSE;

    critical section

    lock := FALSE;
    remainder section
UNTIL FALSE;
```

Critical Section w/ multiple processes

```
// Use TEST-AND-SET
// Common data structures, initialized to false
var waiting: array [0.. n-1] of boolean
    lock: boolean

// For process Pi:
var j: 0..n-1;
    key: boolean;

repeat
    entry section

    waiting[i] := true;
    key := true;
    while waiting[i] and key do key := TEST-AND-SET(lock);
    waiting[i] := false;

    critical section

    j := i+1 mod n;
    while (j != i) and (not waiting[j]) do j := j+1 mod n;
    if j = i then lock := false;
        else waiting[j] := false;

    remainder section
until false;
```

Semaphores

Semaphore S \equiv an integer variable
only accessible through two
atomic operations **P** and **V**

```
P(S): while S <= 0 do skip;  
      S := S - 1;
```

```
V(S): S := S + 1;
```

Example: Let n processes share **lock**

```
// Let lock be initialized to 1  
// Then code for Pi →  
  
repeat  
    P(lock);  
    critical section  
    V(lock);  
    remainder section  
until false;  
  
// this is similar to  
// test-and-set
```

Ex: Semaphores and Synchronization

```
sync := 0;    // not 1
```

P1

.

.

S1;

V(sync);

.

.

P2

.

.


P(sync);

S2;

.

.

// what are the P and V doing?



var a, b, c, d, e, f, g: semaphores;

(* initial value of all semaphores is 0 *)

all in
"wait"
state

begin

parbegin

begin S₁; V(a); V(b); end;

begin P(a); S₂; S₄; V(c); V(d); end;

begin P(b); S₃; V(e); end;

begin P(c); S₅; V(f); end;

begin P(d); ~~P(e)~~; S₆; V(g); end;

begin P(f); P(g); S₇; end;

parend;

end;

modify?

Semaphore Implementation

Problem with both Peterson's SW solution and semaphore: **Busy Waiting**

Solution: Block if S is not positive

Problem: How do you restart blocked process?

Solution: Through the execution of a **V** op. by another process

```
// Semaphore has not only an integer,  
// but also a list of processes
```

```
type Semaphore = record  
    value = integer;  
    L = list of processes;  
end;
```

```
// If P finds integer not positive,  
// then adds self to list and blocks
```

```
P(s):  S.value := S.value - 1;  
        if S.value < 0  
            then begin  
                add process to S.L;  
                block;  
            end;
```

```
// If V finds integer negative, then it knows  
// that another process is waiting  
// Note: lock not necessarily opened after  
// V(S) completion since S.value can still  
// be <= 0
```

```
V(S):  S.value := S.value + 1;  
        if S.value <= 0  
            then begin  
                remove a process P  
                    from S.L;  
                wake up (P);  
            end;
```

Notes

- Integer can be negative in this implementation:
|value| = # of processes waiting
- Linked list can be list of PCB pointers
 - but then must be in system space
- List should be FIFO or some priority strategy (not stack!)
- P & V must remain atomic
 - lock out interrupts (for a uniprocessor)
 - hardware support or software solution to critical section problem (for a multiprocessor)
- But, still busy waiting on entry into semaphore!
 - However, critical section is very short
 - And requires well thought-out system support