

CUDA/GPU – In Depth: misc.

Outline

1. Atomic operations and synchronization
2. Consistency
3. Asynchronous Data Transfers

Example: Shuffling Data

// Reorder values based on keys

// Each thread moves one element

`__global__ void`

```
shuffle(int* prev_array, int* new_array, int* indices)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    new_array[i] = prev_array[indices[i]];
}
```

Host Code

```
int main()
```

```
{
```

// Run grid of N/256 blocks of 256 threads each

```
shuffle<<< N/256, 256>>>(d_old, d_new, d_ind);
```

```
}
```

Blocks must be independent

- Any possible interleaving of blocks should be valid
 - presumed to run to completion without preemption
 - can run in any order
 - can run concurrently OR sequentially
- Blocks may coordinate but not synchronize
 - shared queue pointer: **OK**
 - shared lock: **BAD** ... can easily deadlock
- Independence requirement gives *scalability*

ATOMICS

THE PROBLEM

How do you do global communication?

Finish a grid and start a new one

Global Communication

Finish a kernel and start a new one

All writes from all threads complete before a kernel finishes

```
step1<<<grid1,blk1>>>( ... );  
// The system ensures that all  
// writes from step1 complete.  
step2<<<grid2,blk2>>>( ... );
```

Would need to decompose kernels into “before” and “after” parts

Or, write to a predefined memory location

Race condition! Updates can be lost

Race Conditions

threadId:0

// vector[0] was equal to 0

vector[0] += 5;

...

a = vector[0];

threadId:1917

vector[0] += 1;

...

a = vector[0];

- What is the value of **a** in thread 0?
- What is the value of **a** in thread 1917?
- Thread 0 could have finished execution before 1917 started
- Or the other way around
- Or both are executing at the same time

Answer: not defined by the programming model, can be arbitrary

Atomics

- CUDA provides `atomic` operations to deal with this problem
- An atomic operation guarantees that only a single thread has access to a **piece** of memory while an operation completes
- The name atomic comes from the fact that it is uninterruptable
- No dropped data, but ordering is still arbitrary
- Different types of atomic instructions
 - `atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor}`
- More types in Fermi
- *Atomics are slower than normal load/store*
- *You can have the whole machine queuing on a single location in memory*

Example: Histogram

*// Determine frequency of colors in a picture colors have
// already been converted into ints. Each thread looks
// at one pixel and increments a counter atomically.*

__global__

```
void histogram(int* color, int* buckets)
{
    int i = threadIdx.x
           + blockDim.x * blockIdx.x;
    int c = colors[i];
    atomicAdd(&buckets[c], 1);
}
```

Example: Workqueue

*// For algorithms where the amount of work per item
// is highly non-uniform, it often makes sense
// to continuously grab work from a queue*

__global__

```
void workq(int* work_q, int* q_counter,  
          int* output, int queue_max)  
{  
    int i = threadIdx.x  
        + blockDim.x * blockIdx.x;  
    int q_index =  
        atomicInc(q_counter, queue_max);  
    int result = do_work(work_q[q_index]);  
    output[i] = result;  
}
```


Example: Global Min/Max (Naive)

*// If you require the maximum across all threads
// in a grid, you could do it with a single global
// maximum value, but it will be VERY slow*

`__global__`

```
void global_max(int* values, int* gl_max)
{
    int i = threadIdx.x
           + blockDim.x * blockIdx.x;
    int val = values[i];
    atomicMax(gl_max, val);
}
```

Example: Global Min/Max (Better)

```
// introduce intermediate maximum results, so that  
// most threads do not try to update the global max  
__global__  
void global_max(int* values, int* max,  
                int *regional_maxes,  
                int num_regions)  
{  
    // i and val as before ...  
    int region = i % num_regions;  
    if(atomicMax(&reg_max[region],val) < val)  
    {  
        atomicMax(max,val);  
    }  
}
```

Global Min/Max

- Single value causes serial bottleneck
- Create hierarchy of values for more parallelism
- Performance will still be slow, so use judiciously

Previous solution (based on reduction) much better!

ATOMICS – Summary (so far)

- Can't use normal load/store for inter-thread communication because of **race conditions**
- Use atomic instructions for sparse and/or unpredictable global communication
 - Use shared memory and scan for other communication patterns
- Decompose data for more parallelism
 - But still must have very limited use of single global sum/max/min/etc.

Communication Through Memory

Question:

```
__global__ void race(void)
{
    __shared__ int my_shared_variable;
    my_shared_variable = threadIdx.x;

    // what is the value of
    // my_shared_variable?
}
```

Communication Through Memory

- This is a **race condition**
- The result is **undefined**
- The order in which threads access the variable is undefined without explicit coordination
- Use barriers (e.g., `__syncthreads`) or atomic operations (e.g., `atomicAdd`) to enforce well-defined semantics

Communication Through Memory

```
// Use __syncthreads to ensure data is  
// ready for access  
__global__ void share_data(int *input)  
{  
    __shared__ int data[BLOCK_SIZE];  
  
    data[threadIdx.x] = input[threadIdx.x];  
    __syncthreads();  
  
    // the state of the entire data array is now  
    // well-defined for all threads in this block  
}
```

Communication Through Memory

```
// Use atomic operations to ensure exclusive  
// access to a variable
```

```
// assume *result is initialized to 0
```

```
__global__ void sum(int *input, int *result)
```

```
{
```

```
    atomicAdd(result, input[threadIdx.x]);
```

```
// after this kernel exits, the value of
```

```
// *result will be the sum of the input
```


```
}
```


Resource Contention

- Atomic operations aren't cheap!
- They imply (force) **serialized access** to a variable

```
__global__ void sum(int *input, int *result)
{
    atomicAdd(result, input[threadIdx.x]);
}
...
// how many threads will contend
// for exclusive access to result?
sum<<<B,N/B>>>(input,result);
```

Hierarchical Atomics



- Divide & Conquer
 - Per-thread `atomicAdd` to a `__shared__` partial sum
 - Per-block `atomicAdd` to the total sum

Hierarchical Atomics

```
__global__ void sum(int *input, int *result)
{
    __shared__ int partial_sum;    // shared!

    // thread 0 is responsible for
    // initializing partial_sum
    if(threadIdx.x == 0) partial_sum = 0;
    __syncthreads();

    ...
}
```

Hierarchical Atomics

```
__global__ void sum(int *input, int *result)
{
    ...
    // each thread updates the partial sum
    atomicAdd(&partial_sum,      // shared!
              input[threadIdx.x]);
    __syncthreads();

    // thread 0 updates the total sum
    if(threadIdx.x == 0)
        atomicAdd(result, partial_sum);
}
```

Advice

- Use barriers such as `__syncthreads` to wait until `__shared__` data is ready
- Prefer barriers to atomics when data access patterns are regular or predictable
- Prefer atomics to barriers when data access patterns are sparse or unpredictable
- Atomics with `__shared__` variables are much faster than atomics to `__global__` variables
- Don't synchronize or serialize unnecessarily

Ex: Count 6s

- Global, device functions and excerpts from host, main

```
__device__ int compare(int a, int b) {  
    if (a == b) return 1;  
    return 0;  
}
```

```
__global__ void compute(int *d_in, int *d_out) {  
    d_out[threadIdx.x] =  
    for (i=0; i<SIZE/BLOCKSIZE; i++) {  
        int val = d_in[i*BLOCKSIZE + threadIdx.x];  
        d_out[threadIdx.x] +=  
            compare(val, 6);  
    }  
}
```

*Compute individual
results for each thread
Serialize final results
gathering on host*

```
int __host__ void outer_compute  
    (int *h_in_array, int *h_out_array) {  
    ...  
    compute<<<1, BLOCKSIZE, msize>>>  
        (d_in_array, d_out_array);  
    cudaMemcpy(h_out_array, d_out_array,  
        BLOCKSIZE*sizeof(int),  
        cudaMemcpyDeviceToHost);  
    main(int argc, char **argv) {  
        ...  
        for (int i=0; i<BLOCKSIZE; i++)  
            { sum+=out_array[i]; }  
        printf ("Result = %d\n", sum);  
    }
```

What if we computed sum on GPU?

- Global, device functions and excerpts from host, main

```
__device__ int compare(int a, int b) {  
    if (a == b) return 1;  
    return 0;  
}  
  
__global__ void compute(int *d_in, int *sum) {  
    *sum = 0;  
    for (i=0; i<SIZE/BLOCKSIZE; i++) {  
        int val = d_in[i*BLOCKSIZE + threadIdx.x];  
        *sum += compare(val, 6);  
    }  
}
```

```
int __host__ void outer_compute  
(int *h_in_array, int *h_sum) {  
    ...  
    compute<<<1,BLOCKSIZE,msize>>>  
        (d_in_array, d_sum);  
    cudaThreadSynchronize();  
    cudaMemcpy(h_sum, d_sum,  
        sizeof(int),  
        cudaMemcpyDeviceToHost);  
}
```

```
main(int argc, char **
```

```
...
```

```
    int *sum; // an integer  
    outer_compute(in_array, sum);  
    printf ("Result = %d\n",sum);  
}
```

*Each thread increments
"sum" variable*

Threads Access the Same Memory!

- Global memory and shared memory within an SM can be freely accessed by multiple threads
- Requires appropriate sequencing of memory accesses across threads to same location ***if at least one access is a write***

Using Data Dependences to Reason about Race Conditions

- Compiler research on data dependence analysis provides a systematic way to conservatively identify race conditions on scalar and array variables
 - “Forall” if no dependences cross the iteration boundary of a parallel loop.
 - No loop-carried dependences
 - If a race condition is found,
 - EITHER serialize loop(s) carrying dependence by making it internal to thread program, or part of the host code
 - OR add “synchronization”

Back to our Example: What if Threads Need to Access Same Memory Location

- Dependence on sum across iterations/threads
 - But reordering ok since operations on sum are associative
- Load/increment/store must be done **atomically** to preserve sequential meaning
- Add Synchronization
 - Protect memory locations
 - Control-based (what are threads doing?)
- Definitions:
 - **Atomicity**: a set of operations is atomic if either they all execute or none executes. Thus, there is no way to see the results of a partial execution.
 - **Mutual exclusion**: at most one thread can execute the code at any time
 - **Barrier**: forces threads to stop and wait until all threads have arrived at some point in code, and typically at the same point

Gathering Results on GPU: Barrier Synchronization w/in Block

`void __syncthreads();`

- **Functionality:** Synchronizes all threads in a block
 - Each thread waits at the point of this call until all other threads have reached it
 - Once all threads have reached this point, execution resumes normally
- Why is this needed?
 - A thread can freely read the shared memory of its thread block or the global memory of either its block or grid.
 - Allows the program to guarantee partial ordering of these accesses to prevent incorrect orderings.
- Watch out!
 - Potential for deadlock when it appears in conditionals

Gathering Results on GPU for “Count 6”

```
__global__ void compute(int *d_in, int
    *d_out) {

    d_out[threadIdx.x] = 0;
    for (i=0; i<SIZE/BLOCKSIZE; i++) {
        int val = d_in[i*BLOCKSIZE + threadIdx.x];
        d_out[threadIdx.x] += compare(val, 6);
    }
}
```

```
__global__ void compute(int *d_in, int
    *d_out, int *d_sum) {

    d_out[threadIdx.x] = 0;
    for (i=0; i<SIZE/BLOCKSIZE; i++) {
        int val = d_in[i*BLOCKSIZE + threadIdx.x];
        d_out[threadIdx.x] += compare(val, 6);
    }

    __syncthreads();
    if (threadIdx.x == 0) {
        for 0..BLOCKSIZE-1

            *d_sum += d_out[i];
        }
    }
}
```

Gathering Results on GPU: Atomic Update to Sum Variable

```
int atomicAdd(int* address, int val);
```

Increments the integer at address by val.

Atomic means that once initiated, the operation executes to completion without interruption by other threads

Gathering Results on GPU for “Count 6”

```
__global__ void compute(int *d_in, int
    *d_out) {

    d_out[threadIdx.x] = 0;
    for (i=0; i<SIZE/BLOCKSIZE; i++) {
        int val = d_in[i*BLOCKSIZE + threadIdx.x];
        d_out[threadIdx.x] += compare(val, 6);
    }
}
```

```
__global__ void compute(int *d_in, int
    *d_out, int *d_sum) {

    d_out[threadIdx.x] = 0;
    for (i=0; i<SIZE/BLOCKSIZE; i++) {
        int val = d_in[i*BLOCKSIZE + threadIdx.x];
        d_out[threadIdx.x] += compare(val, 6);
    }

    atomicAdd(d_sum,
        d_out_array[threadIdx.x]);

}
```

Synchronization Within/Across Blocks: Memory Fence Instructions

void __threadfence_block();

- waits until all global and shared memory accesses made by the calling thread prior to call are visible to all threads in the thread block.
- In general, when a thread issues a series of writes to memory in a particular order, other threads may see the effects of these memory writes in a different order.

void __threadfence();

- Similar to above, but visible to all threads in the device for global memory accesses and all threads in the thread block for shared memory accesses.

void __threadfence_system();

- Similar to above, but also visible to host for “page-locked” host memory accesses.

Memory Consistency

Writes to a location become visible to all in the same order

But when does a write become visible

- How to establish orders between a write and a read by different procs?
 - Typically use event synchronization, by using more than one location

| P_1 | P_2 |
|---|----------------------------------|
| /* Assume initial values of A and flag are 0 */ | |
| A = 1; | while (flag == 0); /*spin idly*/ |
| flag = 1; | print A; |

- Intuition not guaranteed by coherence
- Sometimes expect memory to respect order between accesses to *different* locations issued by a given process
 - to preserve orders among accesses to same location by different processes
- Coherence doesn't help: pertains only to single location

Memory Fence Example

```
__device__ unsigned int count = 0;
__shared__ bool isLastBlockDone;
__global__ void sum(const float* array,
                    unsigned int N, float* result) {

    // Each block sums a subset of the input array
    float partialSum = calculatePartialSum(array, N);

    if (threadIdx.x == 0) {
        // Thread 0 of each block stores the partial sum
        // to global memory
        result[blockIdx.x] = partialSum;

        // Thread 0 makes sure its partial sum is written
        // all other threads
        __threadfence();

        // Thread 0 of each block signals that it is done
        unsigned int value = atomicInc(&count, gridDim.x);

        // Thread 0 of each block determines if its block is
        // the last block to be done
        isLastBlockDone = (value == (gridDim.x - 1));
    }
}
```

*Make sure write
to result complete
before continuing*

```
// Synchronize to make sure that each thread
// reads the correct value of isLastBlockDone
__syncthreads();

if (isLastBlockDone) {
    // The last block sums the partial sums
    // stored in result[0 .. gridDim.x-1]
    float totalSum = calculateTotalSum(result);

    (threadIdx.x == 0) {
        // Thread 0 of last block stores total sum
        // to global memory and resets count so
        // that next kernel call works properly
        result[0] = totalSum;
        count = 0;
    }
}
```

Host-Device Transfers (implicit in synchronization discussion)

- Host-Device Data Transfers
 - Device to host memory bandwidth much lower than device to device bandwidth
 - 8 GB/s peak (PCI-e x16 Gen 2) vs. 102 GB/s peak (Tesla C1060)
- Minimize transfers
 - Intermediate data can be allocated, operated on, and deallocated without ever copying them to host memory
- Group transfers
 - One large transfer much better than many small ones

Slide source: Nvidia, 2008

Asynchronous Copy To/From Host

- Warning: I have not tried this!
- Concept:
 - Memory bandwidth can be a limiting factor on GPUs
 - Sometimes computation cost dominated by copy cost
 - But for some computations, data can be “tiled” and computation of tiles can proceed in parallel
 - Can we be computing on one tile while copying another?
- Strategy:
 - Use page-locked memory on host, and asynchronous copies
 - Primitive **cudaMemcpyAsync**
 - Effect is GPU performs DMA from Host Memory
 - Synchronize with **cudaThreadSynchronize()**

Page-Locked Host Memory

- How the Async copy works:
 - DMA performed by GPU memory controller
 - CUDA driver takes virtual addresses and translates them to physical addresses
 - Then copies physical addresses onto GPU
 - Now what happens if the host OS decides to swap out the page???
- Special malloc holds page in place on host
 - Prevents host OS from moving the page
 - `CudaMallocHost()`
- But performance could degrade if this is done on lots of pages!
 - Bypassing virtual memory mechanisms

Example of Asynchronous Data Transfer

```
cudaStreamCreate(&stream1);  
cudaStreamCreate(&stream2);  
cudaMemcpyAsync(dst1, src1, size, dir, stream1);  
kernel<<<grid, block, 0, stream1>>>(...);  
cudaMemcpyAsync(dst2, src2, size, dir, stream2);  
kernel<<<grid, block, 0, stream2>>>(...);
```

*src1 and src2 must have been allocated using cudaMallocHost
stream1 and stream2 identify streams associated with asynchronous
call (note 4th “parameter” to kernel invocation)*

Code from asyncAPI SDK project

```
// allocate host memory
CUDA_SAFE_CALL( cudaMallocHost((void**)&a, nbytes) );
memset(a, 0, nbytes);

// allocate device memory
CUDA_SAFE_CALL( cudaMalloc((void**)&d_a, nbytes) );
CUDA_SAFE_CALL( cudaMemset(d_a, 255, nbytes) );

... // declare grid and thread dimensions and create start and stop events

// asynchronously issue work to the GPU (all to stream 0)
cudaEventRecord(start, 0);
cudaMemcpyAsync(d_a, a, nbytes, cudaMemcpyHostToDevice, 0);
increment_kernel<<<blocks, threads, 0, 0>>>(d_a, value);
cudaMemcpyAsync(a, d_a, nbytes, cudaMemcpyDeviceToHost, 0);
cudaEventRecord(stop, 0);

// have CPU do some work while waiting for GPU to finish

// release resources
CUDA_SAFE_CALL( cudaFreeHost(a) );
CUDA_SAFE_CALL( cudaFree(d_a) );
```