Tutorials | Exercises | Abstracts | LC Workshops | Comments | Search | Privacy & Legal Notice

POSIX Threads Programming

Blaise Barney, Lawrence Livermore National Laboratory

UCRL-MI-133316

Table of Contents

- 1. Abstract
- 2. Pthreads Overview
 - 1. What is a Thread?
 - 2. What are Pthreads?
 - 3. Why Pthreads?
 - 4. <u>Designing Threaded Programs</u>
- 3. The Pthreads API
- 4. Compiling Threaded Programs
- 5. Thread Management
 - 1. Creating and Terminating Threads
 - 2. Passing Arguments to Threads
 - 3. Joining and Detaching Threads
 - 4. Stack Management
 - 5. Miscellaneous Routines
- 6. Mutex Variables
 - 1. Mutex Variables Overview
 - 2. Creating and Destroying Mutexes
 - 3. Locking and Unlocking Mutexes
- 7. Condition Variables
 - 1. Condition Variables Overview
 - 2. Creating and Destroying Condition Variables
 - 3. Waiting and Signaling on Condition Variables
- 8. LLNL Specific Information and Recommendations
- 9. Topics Not Covered
- 10. Pthread Library Routines Reference
- 11. References and More Information
- 12. Exercise

Abstract

In shared memory multiprocessor architectures, such as SMPs, threads can be used to implement parallelism. Historically, hardware vendors have implemented their own proprietary versions of threads, making portability a concern for software developers. For UNIX systems, a standardized C language threads programming interface has been specified by the IEEE POSIX 1003.1c standard. Implementations that adhere to this standard are referred to as POSIX threads, or Pthreads.

The tutorial begins with an introduction to concepts, motivations, and design considerations for using Pthreads. Each of the three major classes of routines in the Pthreads API are then covered: Thread Management, Mutex Variables, and Condition Variables. Example codes are used throughout to demonstrate how to use most of the Pthreads routines needed by a new Pthreads programmer. The tutorial concludes with a discussion of LLNL specifics and how to mix MPI with pthreads. A lab exercise, with numerous example codes (C Language) is also included.

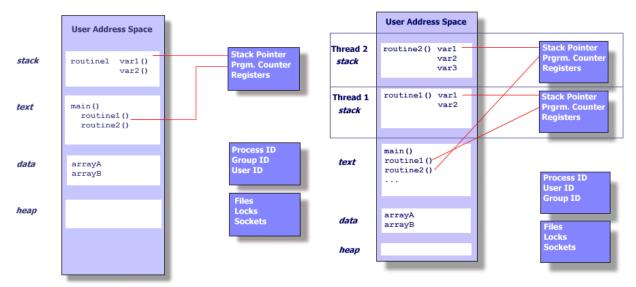
Level/Prerequisites: Ideal for those who are new to parallel programming with threads. A basic understanding of parallel programming in C is assumed. For those who are unfamiliar with Parallel Programming in general, the material covered in EC3500: Introduction To Parallel Computing would be helpful.

Pthreads Overview

What is a Thread?

- Technically, a thread is defined as an independent stream of instructions that can be scheduled to run as such by the operating system. But what
 does this mean?
- To the software developer, the concept of a "procedure" that runs independently from its main program may best describe a thread.
- To go one step further, imagine a main program (a.out) that contains a number of procedures. Then imagine all of these procedures being able to be scheduled to run simultaneously and/or independently by the operating system. That would describe a "multi-threaded" program.
- How is this accomplished?

- Before understanding a thread, one first needs to understand a UNIX process. A process is created by the operating system, and requires a fair amount of "overhead". Processes contain information about program resources and program execution state, including:
 - o Process ID, process group ID, user ID, and group ID
 - o Environment
 - o Working directory.
 - o Program instructions
 - o Registers
 - o Stack
 - o Heap
 - o File descriptors
 - o Signal actions
 - o Shared libraries
 - o Inter-process communication tools (such as message queues, pipes, semaphores, or shared memory).



UNIX PROCESS

- THREADS WITHIN A UNIX PROCESS
- Threads use and exist within these process resources, yet are able to be scheduled by the operating system and run as independent entities largely because they duplicate only the bare essential resources that enable them to exist as executable code.
- This independent flow of control is accomplished because a thread maintains its own:
 - Stack pointer
 - o Registers
 - o Scheduling properties (such as policy or priority)
 - o Set of pending and blocked signals
 - o Thread specific data.
- So, in summary, in the UNIX environment a thread:
 - o Exists within a process and uses the process resources
 - $\circ~$ Has its own independent flow of control as long as its parent process exists and the OS supports it
 - o Duplicates only the essential resources it needs to be independently schedulable
 - o May share the process resources with other threads that act equally independently (and dependently)
 - o Dies if the parent process dies or something similar
 - o Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.
- Because threads within the same process share resources:
 - o Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.
 - o Two pointers having the same value point to the same data.
 - o Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer.

Pthreads Overview

What are Pthreads?

- Historically, hardware vendors have implemented their own proprietary versions of threads. These implementations differed substantially from each other making it difficult for programmers to develop portable threaded applications.
- In order to take full advantage of the capabilities provided by threads, a standardized programming interface was required.
 - o For UNIX systems, this interface has been specified by the IEEE POSIX 1003.1c standard (1995).

- o Implementations adhering to this standard are referred to as POSIX threads, or Pthreads.
- o Most hardware vendors now offer Pthreads in addition to their proprietary API's.
- The POSIX standard has continued to evolve and undergo revisions, including the Pthreads specification. The latest version is known as IEEE Std 1003.1, 2004 Edition.
- Some useful links:
 - o POSIX FAQs: www.opengroup.org/austin/papers/posix_faq.html
 - o Download the Standard: www.unix.org/version3/ieee std.html
- Pthreads are defined as a set of C language programming types and procedure calls, implemented with a pthread.h header/include file and a thread library though this library may be part of another library, such as libe, in some implementations.

Pthreads Overview

Why Pthreads?

- The primary motivation for using Pthreads is to realize potential program performance gains.
- When compared to the cost of creating and managing a process, a thread can be created with much less operating system overhead. Managing threads requires fewer system resources than managing processes.

For example, the following table compares timing results for the fork() subroutine and the pthreads_create() subroutine. Timings reflect 50,000 process/thread creations, were performed with the time utility, and units are in seconds, no optimization flags.

Note: don't expect the system and user times to add up to real time, because these are SMP systems with multiple CPUs working on the problem at the same time. At best, these are approximations run on local machines, past and present.

Dietform		fork()		pthread_create()		
Platform	real	user	sys	real	user	sys
AMD 2.3 GHz Opteron (16cpus/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8cpus/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

Source fork vs thread.txt

- All threads within a process share the same address space. Inter-thread communication is more efficient and in many cases, easier to use than inter-process communication.
- Threaded applications offer potential performance gains and practical advantages over non-threaded applications in several other ways:
 - o Overlapping CPU work with I/O: For example, a program may have sections where it is performing a long I/O operation. While one thread is waiting for an I/O system call to complete, CPU intensive work can be performed by other threads.
 - o Priority/real-time scheduling: tasks which are more important can be scheduled to supersede or interrupt lower priority tasks.
 - Asynchronous event handling: tasks which service events of indeterminate frequency and duration can be interleaved. For example, a web server can both transfer data from previous requests and manage the arrival of new requests.
- The primary motivation for considering the use of Pthreads on an SMP architecture is to achieve optimum performance. In particular, if an
 application is using MPI for on-node communications, there is a potential that performance could be greatly improved by using Pthreads for onnode data transfer instead.
- · For example:
 - MPI libraries usually implement on-node task communication via shared memory, which involves at least one memory copy operation (process to process).
 - For Pthreads there is no intermediate memory copy required because threads share the same address space within a single process. There is
 no data transfer, per se. It becomes more of a cache-to-CPU or memory-to-CPU bandwidth (worst case) situation. These speeds are much
 higher.
 - o Some local comparisons are shown below:

Platform	MPI Shared Memory Bandwidth (GB/sec)	Pthreads Worst Case Memory-to-CPU Bandwidth
----------	---	--

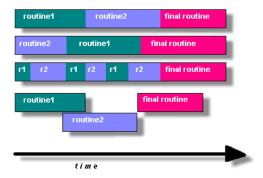
		(GB/sec)
AMD 2.3 GHz Opteron	1.8	5.3
AMD 2.4 GHz Opteron	1.2	5.3
IBM 1.9 GHz POWER5 p5-575	4.1	16
IBM 1.5 GHz POWER4	2.1	4
Intel 2.4 GHz Xeon	0.3	4.3
Intel 1.4 GHz Itanium 2	1.8	6.4

Pthreads Overview

Designing Threaded Programs

Parallel Programming:

- On modern, multi-cpu machines, pthreads are ideally suited for parallel programming, and whatever applies to parallel programming in general, applies to parallel pthreads programs.
- There are many considerations for designing parallel programs, such as:
 - o What type of parallel programming model to use?
 - Problem partitioning
 - Load balancing
 - o Communications
 - o Data dependencies
 - o Synchronization and race conditions
 - o Memory issues
 - o I/O issues
 - o Program complexity
 - o Programmer effort/costs/time
 - o ...
- Covering these topics is beyond the scope of this tutorial, however interested readers can obtain a quick overview in the <u>Introduction to Parallel Computing</u> tutorial.
- In general though, in order for a program to take advantage of Pthreads, it must be able to be organized into discrete, independent tasks which can execute concurrently. For example, if routine1 and routine2 can be interchanged, interleaved and/or overlapped in real time, they are candidates for threading.



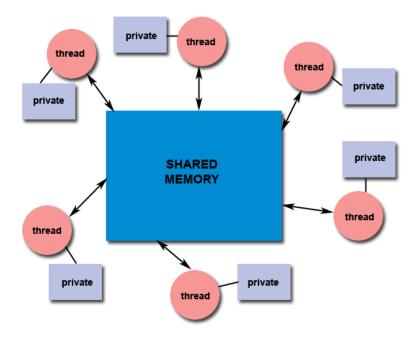
- Programs having the following characteristics may be well suited for pthreads:
 - o Work that can be executed, or data that can be operated on, by multiple tasks simultaneously
 - o Block for potentially long I/O waits
 - o Use many CPU cycles in some places but not others
 - o Must respond to asynchronous events
 - o Some work is more important than other work (priority interrupts)
- Pthreads can also be used for serial applications, to emulate parallel execution. A perfect example is the typical web browser, which for most people, runs on a single cpu desktop/laptop machine. Many things can "appear" to be happening at the same time.
- · Several common models for threaded programs exist:
 - o Manager/worker: a single thread, the manager assigns work to other threads, the workers. Typically, the manager handles all input and

parcels out work to the other tasks. At least two forms of the manager/worker model are common: static worker pool and dynamic worker pool.

- Pipeline: a task is broken into a series of suboperations, each of which is handled in series, but concurrently, by a different thread. An
 automobile assembly line best describes this model.
- o Peer: similar to the manager/worker model, but after the main thread creates other threads, it participates in the work.

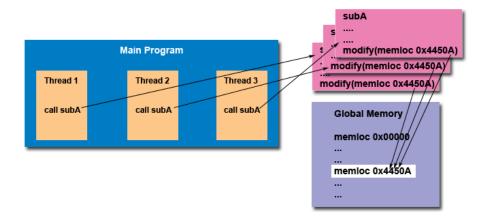
Shared Memory Model:

- · All threads have access to the same global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access (protecting) globally shared data.



Thread-safeness:

- Thread-safeness: in a nutshell, refers an application's ability to execute multiple threads simultaneously without "clobbering" shared data or creating "race" conditions.
- For example, suppose that your application creates several threads, each of which makes a call to the same library routine:
 - o This library routine accesses/modifies a global structure or location in memory.
 - o As each thread calls this routine it is possible that they may try to modify this global structure/memory location at the same time.
 - o If the routine does not employ some sort of synchronization constructs to prevent data corruption, then it is not thread-safe.



• The implication to users of external library routines is that if you aren't 100% certain the routine is thread-safe, then you take your chances with problems that could arise.

• Recommendation: Be careful if your application uses libraries or other objects that don't explicitly guarantee thread-safeness. When in doubt, assume that they are not thread-safe until proven otherwise. This can be done by "serializing" the calls to the uncertain routine, etc.

The Pthreads API

- The original Pthreads API was defined in the ANSI/IEEE POSIX 1003.1 1995 standard. The POSIX standard has continued to evolve and undergo revisions, including the Pthreads specification. The latest version is known as IEEE Std 1003.1, 2004 Edition.
- Copies of the standard can be purchased from IEEE or downloaded for free from www.unix.org/version3/ieee std.html.
- The subroutines which comprise the Pthreads API can be informally grouped into four major groups:
 - 1. **Thread management:** Routines that work directly on threads creating, detaching, joining, etc. They also include functions to set/query thread attributes (joinable, scheduling etc.)
 - Mutexes: Routines that deal with synchronization, called a "mutex", which is an abbreviation for "mutual exclusion". Mutex functions
 provide for creating, destroying, locking and unlocking mutexes. These are supplemented by mutex attribute functions that set or modify
 attributes associated with mutexes.
 - Condition variables: Routines that address communications between threads that share a mutex. Based upon programmer specified conditions. This group includes functions to create, destroy, wait and signal based upon specified variable values. Functions to set/query condition variable attributes are also included.
 - 4. Synchronization: Routines that manage read/write locks and barriers.
- Naming conventions: All identifiers in the threads library begin with **pthread**. Some examples are shown below.

Routine Prefix	Functional Group
pthread_	Threads themselves and miscellaneous subroutines
pthread_attr_	Thread attributes objects
pthread_mutex_	Mutexes
pthread_mutexattr_	Mutex attributes objects.
pthread_cond_	Condition variables
pthread_condattr_	Condition attributes objects
pthread_key_	Thread-specific data keys
pthread_rwlock_	Read/write locks
pthread_barrier_	Synchronization barriers

- The concept of opaque objects pervades the design of the API. The basic calls work to create or modify opaque objects the opaque objects can be modified by calls to attribute functions, which deal with opaque attributes.
- The Pthreads API contains around 100 subroutines. This tutorial will focus on a subset of these specifically, those which are most likely to be immediately useful to the beginning Pthreads programmer.
- For portability, the pthread.h header file should be included in each source file using the Pthreads library.
- The current POSIX standard is defined only for the C language. Fortran programmers can use wrappers around C function calls. Some Fortran compilers (like IBM AIX Fortran) may provide a Fortram pthreads API.
- A number of excellent books about Pthreads are available. Several of these are listed in the References section of this tutorial.

Compiling Threaded Programs

• Several examples of compile commands used for pthreads codes are listed in the table below.

Compiler / Platform	Compiler Command	Description
	xlc_r / cc_r	C (ANSI / non-ANSI)
IBM AIX	xlC_r	C++
	xlf_r -qnosave	Fortran - using IBM's Pthreads API (non-portable)

	xlf90_r -qnosave	
INTEL Linux	icc -pthread	С
	icpc -pthread	C++
PathScale Linux	pathcc -pthread	С
	pathCC -pthread	C++
PGI Linux	pgcc -lpthread	С
	pgCC -lpthread	C++
GNU Linux, AIX	gcc -pthread	GNU C
	g++ -pthread	GNU C++

Thread Management

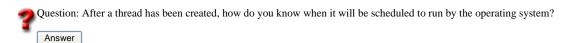
Creating and Terminating Threads

Routines:

```
pthread_create (thread,attr,start_routine,arg)
pthread_exit (status)
pthread_attr_init (attr)
pthread_attr_destroy (attr)
```

Creating Threads:

- Initially, your main() program comprises a single, default thread. All other threads must be explicitly created by the programmer.
- pthread_create creates a new thread and makes it executable. This routine can be called any number of times from anywhere within your code.
- pthread_create arguments:
 - o thread: An opaque, unique identifier for the new thread returned by the subroutine.
 - o attr: An opaque attribute object that may be used to set thread attributes. You can specify a thread attributes object, or NULL for the default values.
 - o start_routine: the C routine that the thread will execute once it is created.
 - o arg: A single argument that may be passed to *start_routine*. It must be passed by reference as a pointer cast of type void. NULL may be used if no argument is to be passed.
- The maximum number of threads that may be created by a process is implementation dependent.
- Once created, threads are peers, and may create other threads. There is no implied hierarchy or dependency between threads.



Thread Attributes:

- By default, a thread is created with certain attributes. Some of these attributes can be changed by the programmer via the thread attribute object.
- $\bullet \ \ \texttt{pthread_attr_init} \ \ \textbf{and} \ \ \texttt{pthread_attr_destroy} \ \ \textbf{are} \ \ \textbf{used} \ \ \textbf{to} \ \ \textbf{initialize/destroy} \ \ \textbf{the} \ \ \textbf{thread} \ \ \textbf{attribute} \ \ \textbf{object}.$
- Other routines are then used to query/set specific attributes in the thread attribute object.
- Some of these attributes will be discussed later.

Terminating Threads:

- There are several ways in which a Pthread may be terminated:
 - \circ The thread returns from its starting routine (the main routine for the initial thread).
 - o The thread makes a call to the pthread_exit subroutine (covered below).

- o The thread is canceled by another thread via the pthread_cancel routine (not covered here).
- o The entire process is terminated due to a call to either the exec or exit subroutines.
- pthread_exit is used to explicitly exit a thread. Typically, the pthread_exit() routine is called after a thread has completed its work and is no longer required to exist.
- If main() finishes before the threads it has created, and exits with pthread_exit(), the other threads will continue to execute. Otherwise, they will be automatically terminated when main() finishes.
- The programmer may optionally specify a termination status, which is stored as a void pointer for any thread that may join the calling thread.
- Cleanup: the pthread_exit() routine does not close files; any files opened inside the thread will remain open after the thread is terminated.
- Discussion: In subroutines that execute to completion normally, you can often dispense with calling pthread_exit() unless, of course, you want to pass a return code back. However, in main(), there is a definite problem if main() completes before the threads it spawned. If you don't call pthread_exit() explicitly, when main() completes, the process (and all threads) will be terminated. By calling pthread_exit() in main(), the process and all of its threads will be kept alive even though all of the code in main() has been executed.

Example: Pthread Creation and Termination

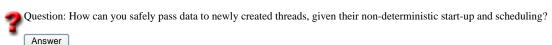
This simple example code creates 5 threads with the pthread_create() routine. Each thread prints a "Hello World!" message, and then terminates with a call to pthread_exit().

```
Example Code - Pthread Creation and Termination
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS
void *PrintHello(void *threadid)
   long tid;
   tid = (long)threadid;
   printf("Hello World! It's me, thread $$\$ld!\n", tid);
   pthread exit(NULL);
int main (int argc, char *argv[])
   pthread_t threads[NUM_THREADS];
   int re;
   long t;
   for(t=0; t<NUM_THREADS; t++){</pre>
      printf("In main: creating thread %ld\n", t);
         = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
         printf("ERROR; return code from pthread_create() is %d\n", rc);
   pthread_exit(NULL);
         Output
```

Thread Management

Passing Arguments to Threads

- The pthread_create() routine permits the programmer to pass one argument to the thread start routine. For cases where multiple arguments must be passed, this limitation is easily overcome by creating a structure which contains all of the arguments, and then passing a pointer to that structure in the pthread_create() routine.
- All arguments must be passed by reference and cast to (void *).



Example 1 - Thread Argument Passing

This code fragment demonstrates how to pass a simple integer to each thread. The calling thread uses a unique data structure for each thread, insuring that each thread's argument remains intact throughout the program.

```
long *taskids[NUM_THREADS];
for(t=0; t<NUM_THREADS; t++)
{
    taskids[t] = (long *) malloc(sizeof(long));
    *taskids[t] = t;
    printf("Creating thread %ld\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) taskids[t]);
    ...
}
</pre>
Source    Output
```

Example 2 - Thread Argument Passing

This example shows how to setup/pass multiple arguments via a structure. Each thread receives a unique instance of the structure.

```
struct thread_data{
  int thread_id;
  int sum;
  char *message;
};
struct thread_data thread_data_array[NUM_THREADS];
void *PrintHello(void *threadarg)
  struct thread_data *my_data;
  my_data = (struct thread_data *) threadarg;
  taskid = my_data->thread_id;
  sum = my data->sum;
  hello_msg = my_data->message;
int main (int argc, char *argv[])
  thread_data_array[t].thread_id = t;
  thread_data_array[t].sum = sum;
  thread_data_array[t].message = messages[t];
  rc = pthread_create(&threads[t], NULL, PrintHello,
       (void *) &thread_data_array[t]);
 Source Output
```

Example 3 - Thread Argument Passing (Incorrect)

This example performs argument passing incorrectly. It passes the *address* of variable t, which is shared memory space and visible to all threads. As the loop iterates, the value of this memory location changes, possibly before the created threads can access it.

```
int rc;
long t;

for(t=0; t<NUM_THREADS; t++)
{
    printf("Creating thread %ld\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &t);
    ...
}

Source    Output</pre>
```

Thread Management

Joining and Detaching Threads

Routines:

```
pthread_join (threadid,status)

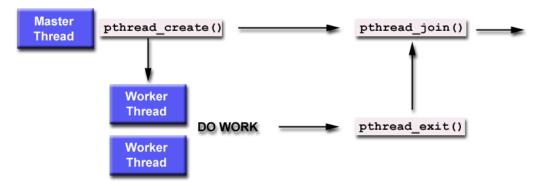
pthread_detach (threadid,status)

pthread_attr_setdetachstate (attr,detachstate)

pthread_attr_getdetachstate (attr,detachstate)
```

Joining:

• "Joining" is one way to accomplish synchronization between threads. For example:



- The pthread_join() subroutine blocks the calling thread until the specified threadid thread terminates.
- The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to pthread_exit().
- A joining thread can match one pthread_join() call. It is a logical error to attempt multiple joins on the same thread.
- Two other synchronization methods, mutexes and condition variables, will be discussed later.

Joinable or Not?

- When a thread is created, one of its attributes defines whether it is joinable or detached. Only threads that are created as joinable can be joined. If a
 thread is created as detached, it can never be joined.
- The final draft of the POSIX standard specifies that threads should be created as joinable.
- To explicitly create a thread as joinable or detached, the attr argument in the pthread_create() routine is used. The typical 4 step process is:
 - 1. Declare a pthread attribute variable of the pthread_attr_t data type
 - 2. Initialize the attribute variable with pthread_attr_init()
 - 3. Set the attribute detached status with pthread_attr_setdetachstate()
 - 4. When done, free library resources used by the attribute with $pthread_attr_destroy()$

Detaching:

- The pthread_detach() routine can be used to explicitly detach a thread even though it was created as joinable.
- There is no converse routine.

Recommendations:

- If a thread requires joining, consider explicitly creating it as joinable. This provides portability as not all implementations may create threads as
 joinable by default.
- If you know in advance that a thread will never need to join with another thread, consider creating it in a detached state. Some system resources may be able to be freed.

Example: Pthread Joining

Example Code - Pthread Joining

This example demonstrates how to "wait" for thread completions by using the Pthread join routine. Since some implementations of Pthreads may not create threads in a joinable state, the threads in this example are explicitly created in a joinable state so that they can be joined later.

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS
void *BusyWork(void *t)
   int i;
   long tid;
   double result=0.0;
   tid = (long)t;
   printf("Thread %ld starting...\n",tid);
   for (i=0; i<1000000; i++)
     result = result + sin(i) * tan(i);
  printf("Thread %ld done. Result = %e\n",tid, result);
  pthread_exit((void*) t);
int main (int argc, char *argv[])
   pthread_t thread[NUM_THREADS];
   pthread_attr_t attr;
   int rc;
  long t;
  void *status;
   /* Initialize and set thread detached attribute */
  pthread_attr_init(&attr);
  pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
   for(t=0; t<NUM_THREADS; t++) {</pre>
     printf("Main: creating thread %ld\n", t);
       c = pthread_create(&thread[t], &attr, BusyWork, (void *)t);
      if (rc) {
        printf("ERROR; return code from pthread_create()
                is %d\n", rc);
         exit(-1);
      }
   /* Free attribute and wait for the other threads */
   for(t=0; t<NUM_THREADS; t++) {
         = pthread_join(thread[t], &status);
      if (rc)
        printf("ERROR; return code from pthread_join()
                is %d\n", rc);
         exit(-1);
     printf("Main: completed join with thread %ld having a status
            of %ld\n",t,(long)status);
printf("Main: program completed. Exiting.\n");
pthread_exit(NULL);
 Source Output
```

Thread Management

Stack Management

Routines:

```
pthread_attr_getstacksize (attr, stacksize)

pthread_attr_setstacksize (attr, stacksize)

pthread_attr_getstackaddr (attr, stackaddr)

pthread_attr_setstackaddr (attr, stackaddr)
```

Preventing Stack Problems:

- The POSIX standard does not dictate the size of a thread's stack. This is implementation dependent and varies.
- Exceeding the default stack limit is often very easy to do, with the usual results: program termination and/or corrupted data.
- Safe and portable programs do not depend upon the default stack limit, but instead, explicitly allocate enough stack for each thread by using the pthread_attr_setstacksize routine.
- The pthread_attr_getstackaddr and pthread_attr_setstackaddr routines can be used by applications in an environment where the stack for a thread must be placed in some particular region of memory.

Some Practical Examples at LC:

• Default thread stack size varies greatly. The maximum size that can be obtained also varies greatly, and may depend upon the number of threads per node.

Node Architecture	#CPUs	Memory (GB)	Default Size (bytes)
AMD Opteron	8	16	2,097,152
Intel IA64	4	8	33,554,432
Intel IA32	2	4	2,097,152
IBM Power5	8	32	196,608
IBM Power4	8	16	196,608
IBM Power3	16	16	98,304

Example: Stack Management

Example Code - Stack Management

This example demonstrates how to query and set a thread's stack size.

```
#include <pthread.h>
#include <stdio.h>
#define NTHREADS 4
#define N 1000
#define MEGEXTRA 1000000
pthread_attr_t attr;
void *dowork(void *threadid)
  double A[N][N];
   int i,j;
   long tid;
  size_t mystacksize;
   tid = (long)threadid;
  pthread_attr_getstacksize (&attr, &mystacksize);
   printf("Thread %ld: stack size = %li bytes \n", tid, mystacksize);
   for (i=0; i< N; i++)
    for (j=0; j< N; j++)

A[i][j] = ((i*j)/3.452) + (N-i);
  pthread_exit(NULL);
int main(int argc, char *argv[])
  pthread_t threads[NTHREADS];
   size_t stacksize;
   int rc;
   long t;
```

```
pthread_attr_init(&attr);
pthread_attr_getstacksize (&attr, &stacksize);
printf("Default stack size = %li\n", stacksize);
stacksize = sizeof(double)*N*N+MEGEXTRA;
printf("Amount of stack needed per thread = %li\n",stacksize);
pthread_attr_setstacksize (&attr, stacksize);
printf("Creating threads with stack size = %li bytes\n",stacksize);
for(t=0; t<NTHREADS; t++){
    rc = pthread_create(&threads[t], &attr, dowork, (void *)t);
    if (rc){
        printf("ERROR; return code from pthread_create() is %d\n", rc);
        exit(-1);
    }
}
printf("Created %ld threads.\n", t);
pthread_exit(NULL);</pre>
```

Thread Management

Miscellaneous Routines

```
pthread self ()
pthread equal (thread1,thread2)
```

- pthread_self returns the unique, system assigned thread ID of the calling thread.
- pthread_equal compares two thread IDs. If the two IDs are different 0 is returned, otherwise a non-zero value is returned.
- Note that for both of these routines, the thread identifier objects are opaque and can not be easily inspected. Because thread IDs are opaque objects, the C language equivalence operator == should not be used to compare two thread IDs against each other, or to compare a single thread ID against another value.

```
pthread_once (once_control, init_routine)
```

- pthread_once executes the init_routine exactly once in a process. The first call to this routine by any thread in the process executes the given init_routine, without parameters. Any subsequent call will have no effect.
- The init_routine routine is typically an initialization routine.
- The once_control parameter is a synchronization control structure that requires initialization prior to calling pthread_once. For example:

```
pthread_once_t once_control = PTHREAD_ONCE_INIT;
```

Mutex Variables

Overview

- Mutex is an abbreviation for "mutual exclusion". Mutex variables are one of the primary means of implementing thread synchronization and for
 protecting shared data when multiple writes occur.
- A mutex variable acts like a "lock" protecting access to a shared data resource. The basic concept of a mutex as used in Pthreads is that only one thread can lock (or own) a mutex variable at any given time. Thus, even if several threads try to lock a mutex only one thread will be successful. No other thread can own that mutex until the owning thread unlocks that mutex. Threads must "take turns" accessing protected data.
- Mutexes can be used to prevent "race" conditions. An example of a race condition involving a bank transaction is shown below:

Thread 1	Thread 2	Balance
Read balance: \$1000		\$1000
	Read balance: \$1000	\$1000
	Deposit \$200	\$1000
Deposit \$200		\$1000

Update balance \$1000+\$200		\$1200
	Update balance \$1000+\$200	\$1200

- In the above example, a mutex should be used to lock the "Balance" while a thread is using this shared data resource.
- Very often the action performed by a thread owning a mutex is the updating of global variables. This is a safe way to ensure that when several
 threads update the same variable, the final value is the same as what it would be if only one thread performed the update. The variables being
 updated belong to a "critical section".
- A typical sequence in the use of a mutex is as follows:
 - o Create and initialize a mutex variable
 - Several threads attempt to lock the mutex
 - o Only one succeeds and that thread owns the mutex
 - o The owner thread performs some set of actions
 - o The owner unlocks the mutex
 - o Another thread acquires the mutex and repeats the process
 - o Finally the mutex is destroyed
- . When several threads compete for a mutex, the losers block at that call an unblocking call is available with "trylock" instead of the "lock" call.
- When protecting shared data, it is the programmer's responsibility to make sure every thread that needs to use a mutex does so. For example, if 4
 threads are updating the same data, but only one uses a mutex, the data can still be corrupted.

Mutex Variables

Creating and Destroying Mutexes

Routines:

```
pthread mutex init (mutex,attr)
pthread mutex destroy (mutex)
pthread mutexattr init (attr)
pthread mutexattr destroy (attr)
```

Usage:

- Mutex variables must be declared with type pthread_mutex_t, and must be initialized before they can be used. There are two ways to initialize a mutex variable:
 - Statically, when it is declared. For example: pthread_mutex_t mymutex = PTHREAD_MUTEX_INITIALIZER;
 - 2. Dynamically, with the pthread_mutex_init() routine. This method permits setting mutex object attributes, attr.

The mutex is initially unlocked.

- The *attr* object is used to establish properties for the mutex object, and must be of type pthread_mutexattr_t if used (may be specified as NULL to accept defaults). The Pthreads standard defines three optional mutex attributes:
 - o Protocol: Specifies the protocol used to prevent priority inversions for a mutex.
 - o Prioceiling: Specifies the priority ceiling of a mutex.
 - o Process-shared: Specifies the process sharing of a mutex.

Note that not all implementations may provide the three optional mutex attributes.

- The pthread_mutexattr_init() and pthread_mutexattr_destroy() routines are used to create and destroy mutex attribute objects respectively.
- pthread_mutex_destroy() should be used to free a mutex object which is no longer needed.

Mutex Variables

Locking and Unlocking Mutexes

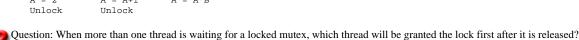
Routines:

```
pthread mutex lock (mutex)
pthread mutex trylock (mutex)
pthread mutex unlock (mutex)
```

Usage:

- The pthread_mutex_lock() routine is used by a thread to acquire a lock on the specified *mutex* variable. If the mutex is already locked by another thread, this call will block the calling thread until the mutex is unlocked.
- pthread_mutex_trylock() will attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with a "busy" error code. This routine may be useful in preventing deadlock conditions, as in a priority-inversion situation.
- pthread_mutex_unlock() will unlock a mutex if called by the owning thread. Calling this routine is required after a thread has completed its use of protected data if other threads are to acquire the mutex for their work with the protected data. An error will be returned if:
 - o If the mutex was already unlocked
 - o If the mutex is owned by another thread
- There is nothing "magical" about mutexes...in fact they are akin to a "gentlemen's agreement" between participating threads. It is up to the code writer to insure that the necessary threads all make the mutex lock and unlock calls correctly. The following scenario demonstrates a logical error:

```
Thread 1 Thread 2 Thread 3
Lock Lock
A = 2 A = A+1 A = A*B
Unlock Unlock
```



Answer

Example: Using Mutexes

Example Code - Using Mutexes

This example program illustrates the use of mutex variables in a threads program that performs a dot product. The main data is made available to all threads through a globally accessible structure. Each thread works on a different part of the data. The main thread waits for all the threads to complete their computations, and then it prints the resulting sum.

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
The following structure contains the necessary information
to allow the function "dotprod" to access its input data and
place its output into the structure. */
typedef struct
   double
              *b;
   double
   double
              sum;
          veclen;
 } DOTDATA;
/* Define globally accessible variables and a mutex */
#define NUMTHRDS 4
#define VECLEN 100
  DOTDATA dotstr;
   pthread_t callThd[NUMTHRDS];
   pthread_mutex_t mutexsum;
The function dotprod is activated when the thread is created.
All input to this routine is obtained from a structure
```

```
of type DOTDATA and all output from this function is written into this structure. The benefit of this approach is apparent for the \,
multi-threaded program: when a thread is created we pass a single
argument to the activated function - typically this argument
is a thread number. All the other information required by the
function is accessed from the globally accessible structure.
void *dotprod(void *arg)
   /* Define and use local variables for convenience */
   int i, start, end, len;
   long offset;
   double mysum, *x, *y;
offset = (long)arg;
   len = dotstr.veclen;
   start = offset*len;
   end = start + len;
   x = dotstr.a;
   y = dotstr.b;
   Perform the dot product and assign result
   to the appropriate variable in the structure.
   mysum = 0;
   for (i=start; i<end; i++)
      mysum += (x[i] * y[i]);
   Lock a mutex prior to updating the value in the shared
   structure, and unlock it upon updating.
   pthread_mutex_lock (&mutexsum);
   dotstr.sum += mysum;
   pthread_mutex_unlock (&mutexsum);
   pthread_exit((void*) 0);
The main program creates threads which do all the work and then
print out result upon completion. Before creating the threads,
the input data is created. Since all threads update a shared structure,
we need a mutex for mutual exclusion. The main thread needs to wait for
all threads to complete, it waits for each one of the threads. We specify
a thread attribute value that allow the main thread to join with the
threads it creates. Note also that we free up handles when they are
no longer needed.
int main (int argc, char *argv[])
   long i;
   double *a, *b;
   void *status;
   pthread_attr_t attr;
   /* Assign storage and initialize values */
   a = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));
   b = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));
   for (i=0; i<VECLEN*NUMTHRDS; i++)</pre>
     a[i]=1.0;
     b[i]=a[i];
   dotstr.veclen = VECLEN;
   dotstr.a = a;
   dotstr.b = b;
   dotstr.sum=0;
   pthread_mutex_init(&mutexsum, NULL);
   /* Create threads to perform the dotproduct */
   pthread_attr_init(&attr);
   pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
        for(i=0; i<NUMTHRDS; i++)</pre>
```

```
Each thread works on a different set of data.
      The offset is specified by 'i'. The size of
      the data for each thread is indicated by VECLEN.
      pthread_create(&callThd[i], &attr, dotprod, (void *)i);
      pthread_attr_destroy(&attr);
      /* Wait on the other threads */
      for(i=0; i<NUMTHRDS; i++)</pre>
        pthread_join(callThd[i], &status);
 /* After joining, print out the results and cleanup */
 printf ("Sum = %f \n", dotstr.sum);
 free (a);
 free (b);
 pthread_mutex_destroy(&mutexsum);
 pthread_exit(NULL);
Source Serial version
Source Pthreads version
```

Condition Variables

Overview

- Condition variables provide yet another way for threads to synchronize. While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.
- Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met. This can be very resource consuming since the thread would be continuously busy in this activity. A condition variable is a way to achieve the same goal without polling.
- · A condition variable is always used in conjunction with a mutex lock.
- A representative sequence for using condition variables is shown below.

Main Thread

- Declare and initialize global data/variables which require synchronization (such as "count")
- Declare and initialize a condition variable object
- Declare and initialize an associated mutex
- · Create threads A and B to do work

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call pthread_cond_wait() to perform a blocking wait for signal from Thread-B. Note that a call to pthread_cond_wait() automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Main Thread

Join / Continue

Thread B

- Do work
- · Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Condition Variables

Creating and Destroying Condition Variables

Routines:

```
pthread_cond_init (condition,attr)
pthread_cond_destroy (condition)
pthread_condattr_init (attr)
pthread_condattr_destroy (attr)
```

Usage:

- Condition variables must be declared with type pthread_cond_t, and must be initialized before they can be used. There are two ways to initialize a condition variable:
 - Statically, when it is declared. For example: pthread_cond_t myconvar = PTHREAD_COND_INITIALIZER;
 - 2. Dynamically, with the pthread_cond_init() routine. The ID of the created condition variable is returned to the calling thread through the *condition* parameter. This method permits setting condition variable object attributes, *attr*.
- The optional *attr* object is used to set condition variable attributes. There is only one attribute defined for condition variables: process-shared, which allows the condition variable to be seen by threads in other processes. The attribute object, if used, must be of type pthread_condattr_t (may be specified as NULL to accept defaults).

Note that not all implementations may provide the process-shared attribute.

- The pthread_condattr_init() and pthread_condattr_destroy() routines are used to create and destroy condition variable attribute objects.
- pthread_cond_destroy() should be used to free a condition variable that is no longer needed.

Condition Variables

Waiting and Signaling on Condition Variables

Routines:

```
pthread cond wait (condition, mutex)
pthread cond signal (condition)
pthread cond broadcast (condition)
```

Usage:

- pthread_cond_wait() blocks the calling thread until the specified *condition* is signalled. This routine should be called while *mutex* is locked, and it will automatically release the mutex while it waits. After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread. The programmer is then responsible for unlocking *mutex* when the thread is finished with it.
- The pthread_cond_signal() routine is used to signal (or wake up) another thread which is waiting on the condition variable. It should be called after *mutex* is locked, and must unlock *mutex* in order for pthread_cond_wait() routine to complete.
- The pthread_cond_broadcast() routine should be used instead of pthread_cond_signal() if more than one thread is in a blocking wait state.
- It is a logical error to call $pthread_cond_signal()$ before calling $pthread_cond_wait()$.

Proper locking and unlocking of the associated mutex variable is essential when using these routines. For example:

- \bullet Failing to lock the mutex before calling ${\tt pthread_cond_wait()} \ may \ cause \ it \ NOT \ to \ block.$
- Failing to unlock the mutex after calling pthread_cond_signal() may not allow a matching pthread_cond_wait() routine to complete (it will remain blocked).

Example: Using Condition Variables

Example Code - Using Condition Variables

This simple example code demonstrates the use of several Pthread condition variable routines. The main routine creates three threads. Two of the threads perform work and update a "count" variable. The third thread waits until the count variable reaches a specified value.

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 3
#define TCOUNT 10
#define COUNT_LIMIT 12
       count = 0;
int
       thread_ids[3] = \{0,1,2\};
int
pthread_mutex_t count_mutex;
pthread_cond_t count_threshold_cv;
void *inc_count(void *t)
  int i;
  long my_id = (long)t;
  for (i=0; i<TCOUNT; i++) {
    pthread_mutex_lock(&count_mutex);
    count++;
    Check the value of count and signal waiting thread when condition is
    reached. Note that this occurs while mutex is locked.
    if (count == COUNT_LIMIT) {
      pthread_cond_signal(&count_threshold_cv);
      printf("inc_count(): thread %ld, count = %d Threshold reached.\n",
             my_id, count);
    printf("inc_count(): thread %ld, count = %d, unlocking mutex\n",
          my_id, count);
    pthread_mutex_unlock(&count_mutex);
    /* Do some "work" so threads can alternate on mutex lock */
    sleep(1);
  pthread_exit(NULL);
void *watch_count(void *t)
  long my_id = (long)t;
  printf("Starting watch_count(): thread %ld\n", my_id);
  Lock mutex and wait for signal. Note that the pthread_cond_wait
  routine will automatically and atomically unlock mutex while it waits.
  Also, note that if COUNT_LIMIT is reached before this routine is run by
  the waiting thread, the loop will be skipped to prevent pthread_cond_wait
  from never returning.
  pthread_mutex_lock(&count_mutex);
  if (count<COUNT_LIMIT) {
    pthread cond wait(&count threshold cv. &count mutex);
    printf("watch_count(): thread %ld Condition signal received.\n", my_id);
    count += 125;
    printf("watch_count(): thread %ld count now = %d.\n", my_id, count);
  pthread_mutex_unlock(&count_mutex);
  pthread_exit(NULL);
int main (int argc, char *argv[])
  int i. rc;
 long t1=1, t2=2, t3=3; pthread_t threads[3];
  pthread attr t attr;
  /* Initialize mutex and condition variable objects */
 pthread_mutex_init(&count_mutex, NULL);
  pthread_cond_init (&count_threshold_cv, NULL);
  /* For portability, explicitly create threads in a joinable state */
  pthread_attr_init(&attr);
```

LLNL Specific Information and Recommendations

This section describes details specific to Livermore Computing's systems.

Implementations:

- All LC production systems include a Pthreads implementation that follows draft 10 (final) of the POSIX standard. This is the preferred implementation.
- Implementations differ in the maximum number of threads that a process may create. They also differ in the default amount of thread stack space.

Compiling:

- LC maintains a number of compilers, and usually several different versions of each see the LC's Supported Compilers web page.
- The compiler commands described in the Compiling Threaded Programs section apply to LC systems.
- Additionally, all LC IBM compilers are aliased to their thread-safe command. For example, xlc really uses xlc_r. This is only true for LC IBM systems.

Mixing MPI with Pthreads:

- Programs that contain both MPI and Pthreads are common and easy to develop on all LC systems.
- Design:
 - o Each MPI process typically creates and then manages N threads, where N makes the best use of the available CPUs/node.
 - o Finding the best value for N will vary with the platform and your application's characteristics.
 - o For IBM SP systems with two communication adapters per node, it may prove more efficient to use two (or more) MPI tasks per node.
 - o In general, there may be problems if multiple threads make MPI calls. The program may fail or behave unexpectedly. If MPI calls must be made from within a thread, they should be made only by one thread.
- Compiling:
 - o Use the appropriate MPI compile command for the platform and language of choice
 - o Be sure to include the required flag as in the table above (-pthread or -qnosave)
 - MPICH is not thread safe
- An example code that uses both MPI and Pthreads is available below. The serial, threads-only, MPI-only and MPI-with-threads versions
 demonstrate one possible progression.
 - o Serial
 - o Pthreads only
 - o MPI only
 - o MPI with pthreads
 - o makefile (for IBM SP)

Topics Not Covered

Several features of the Pthreads API are not covered in this tutorial. These are listed below. See the Pthread-Library Routines Reference section for more information.

- · Thread Scheduling
 - o Implementations will differ on how threads are scheduled to run. In most cases, the default mechanism is adequate.
 - o The Pthreads API provides routines to explicitly set thread scheduling policies and priorities which may override the default mechanisms.
 - o The API does not require implementations to support these features.
- Keys: Thread-Specific Data
 - o As threads call and return from different routines, the local data on a thread's stack comes and goes.
 - To preserve stack data you can usually pass it as an argument from one routine to the next, or else store the data in a global variable associated with a thread.
 - o Pthreads provides another, possibly more convenient and versatile, way of accomplishing this through keys.
- Mutex Protocol Attributes and Mutex Priority Management for the handling of "priority inversion" problems.
- Condition Variable Sharing across processes
- Thread Cancellation
- · Threads and Signals
- · Synchronization constructs barriers and locks

Pthread Library Routines Reference

For convenience, an alphabetical list of Pthread routines, linked to their corresponding man page, is provided below.

pthread atfork

pthread attr destroy

pthread attr getdetachstate

pthread attr getguardsize

pthread_attr_getinheritsched

pthread_attr_getschedparam

pthread attr getschedpolicy

pthread_attr_getscope

pthread_attr_getstack

pthread_attr_getstackaddr

pthread attr getstacksize

pthread attr init

pthread attr setdetachstate

pthread attr_setguardsize

pthread attr setinheritsched

pthread_attr_setschedparam

pthread_attr_setschedpolicy

pthread_attr_setscope pthread_attr_setstack

pthread_attr_setstackaddr

pthread_attr_setstacksize

pthread barrier destroy

pthread barrier init

pthread barrier wait

pthread barrierattr destroy

pthread barrierattr getpshared

pthread_barrierattr_init

pthread_barrierattr_setpshared

pthread_cancel

pthread_cleanup_pop

pthread_cleanup_push

pthread cond broadcast

pthread cond destroy pthread cond init

pthread cond signal

pthread cond timedwait

pthread_cond_wait

pthread_condattr_destroy

pthread_condattr_getclock

pthread_condattr_getpshared

pthread_condattr_init

pthread condattr setclock

pthread condattr setpshared

pthread_create pthread detach pthread equal pthread exit pthread getconcurrency pthread_getcpuclockid pthread_getschedparam pthread_getspecific pthread_join pthread key create pthread_key_delete pthread_kill pthread_mutex_destroy pthread mutex getprioceiling pthread mutex init pthread mutex lock pthread mutex setprioceiling pthread_mutex_timedlock pthread_mutex_trylock pthread_mutex_unlock pthread_mutexattr_destroy pthread_mutexattr_getprioceiling pthread mutexattr getprotocol pthread mutexattr getpshared pthread mutexattr gettype pthread mutexattr init pthread_mutexattr_setprioceiling pthread_mutexattr_setprotocol pthread_mutexattr_setpshared pthread_mutexattr_settype pthread_once pthread_rwlock_destroy pthread rwlock init pthread rwlock rdlock pthread rwlock timedrdlock pthread rwlock timedwrlock pthread_rwlock_tryrdlock pthread rwlock trywrlock pthread_rwlock_unlock pthread_rwlock_wrlock pthread_rwlockattr_destroy pthread_rwlockattr_getpshared pthread rwlockattr init pthread rwlockattr setpshared pthread self pthread setcancelstate pthread_setcanceltype pthread_setconcurrency pthread_setschedparam pthread_setschedprio pthread_setspecific pthread_sigmask pthread_spin_destroy pthread_spin_init pthread spin lock pthread spin trylock

This completes the tutorial.

pthread spin unlock pthread testcancel



Please complete the online evaluation form - unless you are doing the exercise, in which case please complete it at the end of the exercise.

Where would you like to go now?

- Exercise
- Agenda
- Back to the top

References and More Information

- Author: Blaise Barney, Livermore Computing.
- POSIX Standard: www.unix.org/version3/ieee_std.html
- "Pthreads Programming". B. Nichols et al. O'Reilly and Associates.
- "Threads Primer". B. Lewis and D. Berg. Prentice Hall
- "Programming With POSIX Threads". D. Butenhof. Addison Wesley www.awl.com/cseng/titles/0-201-63392-2
- "Programming With Threads". S. Kleiman et al. Prentice Hall

https://computing.llnl.gov/tutorials/pthreads/ Last Modified: 01/13/2010 18:49:19 blaiseb@llnl.gov UCRL-MI-133316