

# How to Write Fast Code

18-645, spring 2008

1<sup>st</sup> Lecture, Jan. 14<sup>th</sup>

**Instructor:** Markus Püschel

**TAs:** Srinivas Chellappa (Vas) and Frédéric de Mesmay (Fred)

# Today

- Motivation and idea behind this course
- Technicalities
- Motivation: Concrete applications

# Motivation and idea behind this course

# Scope

- **Numerical computing:** algorithms and implementation that are dominated by additions and multiplications, usually floating point
- Three domains of numerical computing:


Domain	Platform	Examples
Scientific computing	Large computer clusters	Climate modeling, Physics simulations
Consumer computing	Standard desktop	Adobe Photoshop, Audio/Video coding
Embedded computing	Small low-power processor	Signal processing, Control

- Usually there is an **unlimited need for performance**  
large datasets, realtime

# The Problem

Discrete Fourier Transform (DFT) on 2 x Core 2 Duo 3 GHz (single precision)

Gflop/s




- Standard desktop computer, vendor compiler, using optimization flags
- All implementations have roughly the same operations count ( $\sim 4n\log_2(n)$ )
- *Maybe the DFT is just difficult?*

# The Problem

## Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)


Gflop/s



- Standard desktop computer, vendor compiler, using optimization flags
- All implementations have **exactly** the same operations count ( $2n^3$ )
- *What is going on?*


# Evolution of Processors (Intel)

**Floating point peak performance [Mflop/s]**  
**CPU frequency [MHz]**




data: [www.sandpile.org](http://www.sandpile.org)

# Evolution of Processors (Intel)



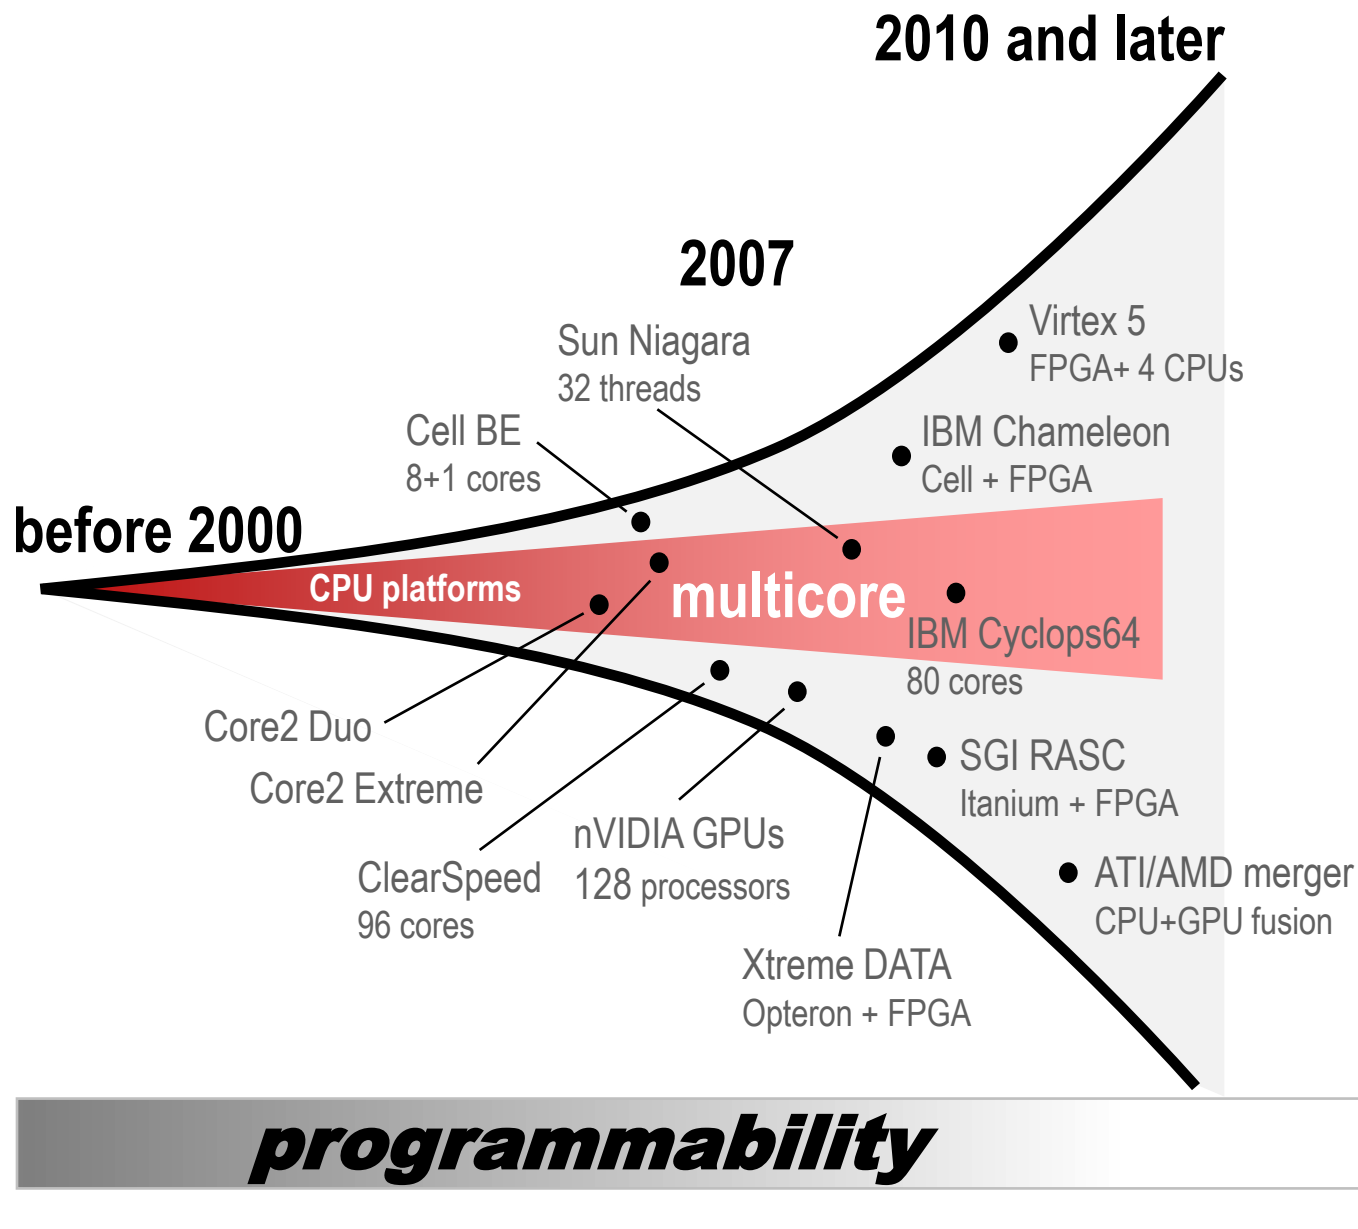


# Evolution of Processors (Intel)



*High performance software development becomes a nightmare*

# Evolution of Processors: The Future




*A clean slate  
for concurrent  
architectures*

# DFT Plot: Analysis


Discrete Fourier Transform (DFT) on 2 x Core 2 Duo 3 GHz

Gflop/s



# MMM Plot: Analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz  
Gflop/s



# Summary and Facts I


- **Implementations with same operations count can have vastly different performance (up to 100x and more)**
  - A cache miss can be 100x more expensive than an addition or multiplication
  - Vector instructions can perform 2 or 3 operations in parallel
  - All recent desktop computers have multiple cores = processors on one die
- **Minimizing operations count does not mean maximizing performance**
- **End of free speed-up: Legacy code will not get automatically faster anymore**
  - CPU frequency scaling has hit the power wall
  - Future performance gains through increasing parallelism
  - It is not clear how future platforms will look

# Summary and Facts II

- **It is very difficult to write the fastest code**
  - Tuning for memory hierarchy
  - Efficient use of vector instructions
  - Efficient parallelization (multiple threads)
  - Requires expert knowledge in algorithms, coding, and architecture
  
- **Compilers can rarely perform the necessary optimization on numerical code**
  - Often intricate changes in the algorithm required
  - Automatic parallelization/vectorization still unsolved
  
- **Highest performance is in general non-portable**
  - Best code on one computer may be suboptimal on another
  - Best code is tuned to microarchitecture
  - Often assembly code is hand-written for optimal tuning

# Current Practice

- **Legions** of programmers implement and optimize the **same** functionality for **every** platform and **whenever** a new platform comes out




# Current Research: Automatic Performance Tuning

- Automate (parts of) the implementation or optimization



- Research efforts

- Linear algebra: **Phipac/ATLAS**, LAPACK, Sparsity/Bebop/OSKI, Flame
- Tensor computations
- PDE/finite elements: Fenics
- Adaptive sorting
- Fourier transform: FFTW
- Linear transforms: Spiral
- ...others
- New compiler techniques






# This Course

- Learn how to write **fast code** for numerical problems
  - Requires multi-disciplinary knowledge
  - Principles studied using important examples
  - Applied in homeworks and a semester-long research project

## Fast implementations of numerical problems



# This Course cont'd

## ■ Background

- Algorithm analysis
- Compilers
- Computer architecture

## ■ Performance optimization

- Benchmarking, optimization techniques (memory hierarchy, vector instructions, multithreading)
- Case studies: important numerical kernels (transforms, linear algebra, filters, convolution, ...)
- Automatic performance tuning (state-of-the-art research)

## ■ Other knowledge

- History, tips for publishing and presenting, ...

# About this Course

## ■ Requirements

- solid C programming skills
- matrix algebra
- senior or above

## ■ Grading

- 40% research project
- 15% midterm
- 35% homework
- 10% class participation

## ■ No textbook

## ■ Office Hours: yet to be determined

## ■ Website: [www.ece.cmu.edu/~pueschel](http://www.ece.cmu.edu/~pueschel) → teaching → 18-645

# Research Project

- Team up in pairs
- **Topic:** Very fast implementation of a numerical problem
- **Jan 28<sup>th</sup>:** suggest to me a problem or I give you a problem  
Tip: pick something from your research (for PhD students)
- Show “milestones” during semester
- Write 4 page standard conference paper (template will be provided)
- Give short presentation end of semester

# Midterm

- Mostly about algorithm analysis
- Some multiple-choice

# Final Exam

- There is no final exam

# Homework

- **Exercises on algorithm analysis (Math)**
- **Implementation exercises**
  - Concrete numerical problems
  - Study the effect of program optimizations, use of compilers, use of special instructions, etc. (Writing C code + creating runtime/performance plots)
  - Some templates will be provided
- **Homework scheduled to leave time for research project**

# Classes/Class Participation

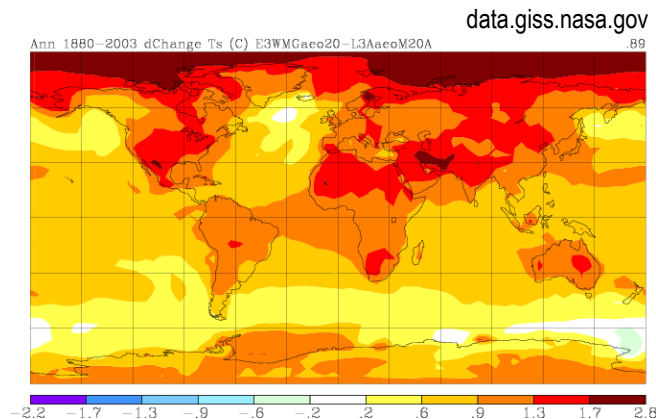
- **I'll start on time, duration ~1:30 (without break)**
  - be on time, it's good style
  
- **It is important to attend**
  - many things I'll teach are not in books
  - I'll use part slides part blackboard
  
- **Ask questions**
  
- **I will provide some anonymous feedback mechanism (maybe every 3-4 weeks)**

# Questions?



# Motivation: Concrete Applications


# Scientific Computing (Large Clusters)



**Climate modelling**



**Finance simulations**



**Molecular dynamics**

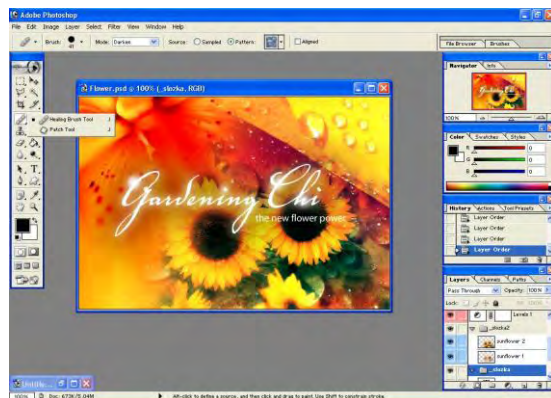
## Other application areas:

- Fluid dynamics
- Chemistry
- Biology
- Medicine
- Geophysics

## Methods:

- Mostly linear algebra
- PDE solving
- Linear system solving
- Finite element methods

# Consumer Computing (Desktop, ...)



Photo/video processing



Audio coding



Security




Image compression

## Methods:

- Linear algebra
- Transforms
- Filters
- Many others


# Embedded Computing (Low-power processors)

[www.dei.unipd.it](http://www.dei.unipd.it)




**Sensor networks**

[www.ece.drexel.edu](http://www.ece.drexel.edu)



**Cars**

[www.microway.com.au](http://www.microway.com.au)



**Robotics**

## Computation needed:

- Signal processing
- Control
- Communication

## Methods:

- Linear algebra
- Transforms, Filters
- Coding


# Research (Examples at ECE/CMU)

Bhagavatula/Savvides




**Biometrics**

Moura



**Medical Imaging**

Kovacevic



**Bioimaging**

Kanade



**Computer vision**

# Summary

- A very large number of diverse applications in engineering, science, consumer market rely on numerical computation
- The computations are diverse but rely on basic mathematical **functionality** (see 13 dwarfs, Berkeley report on parallel computing landscape)
  - Linear algebra (dense/sparse)
  - Transforms/filters
  - Grid methods
  - Encryption
  - Graph traversals, sorting
  - ...
- Unlimited need for performance
- In this course you learn how to make numerical applications run fast on modern computing platforms (focus desktop)